blinkit

PRODUCT CASE STUDY · BLINKIT · QUICK COMMERCE

Why do Blinkit's most loyal users never widen their basket?

Four parts: an AI discovery engine over public reviews, primary user research, a problem definition, and a deployed AI-native MVP — the Category Nudge Agent.

THE CUSTOMER BASE WE ARE DESIGNING INTO

31.8M
avg. monthly transacting customers
16.9M a year earlier

331M
orders in the quarter
≈10.4 orders per customer

2,443
dark stores live
+200 in the quarter

86%
YoY net order value growth
₹17,132 cr NOV

Blinkit segment, Eternal Q1 FY27 (quarter ended 30 Jun 2026) shareholders' letter and earnings call. Growth came from more customers and more orders, not higher basket value — average order value dipped slightly.

DISCOVERY ON BLINKIT TODAY

1. **Intent-led search and category aisles.** AI-assisted search ranking is the primary way a basket gets filled — the customer has to ask first.
2. **Home rails tuned to place and time.** What is popular nearby, time-of-day sets, past purchases, and offer rails on things the customer already looked at.
3. **Complement prompts at cart.** Add-on suggestions next to what is already in the basket — adjacent to intent, rarely a new category.

The gap this study works in: every surface above ranks *what* to buy inside a category the customer already trusts. None of them argue *why* to try a category a bad order taught them to avoid.

Discovery surfaces described from the public app and published reporting, not internal Blinkit documentation.

blinkit | THE GOAL, FRAMED

Heavy repeat buyers keep narrow baskets — and most of them say so themselves

58%

of surveyed users self-report that they mostly stick to the same categories

Source: primary survey, 18 of 31 responses. Survey-based; no live interviews.

This is a trust problem, not an awareness problem

These customers order weekly or daily. They have already seen the other categories in the app. What stops them is what happened — or what they expect to happen — when the order arrives.

The company's own growth path runs through this behaviour

Analyst, Eternal Q4 FY26 earnings call (28 Apr 2026): "a large part of our growth narrative from here on depends in some sense on either growing the non-grocery assortment and going outside of the metro cities."

CFO: "It's a function of assortment expansion, geographical expansion as well as more demand densification in the cities where we are present today."

CAVEAT The call does not discuss product-quality or refund friction. The link to this project's cause is ours, not the company's.

A five-stage pipeline hard-gates on category-adoption relevance before it extracts anything

LIVE · DISCOVERY WORKFLOW
Open the live extractor →

15,820 public reviews ingested → **4,740** pass the rule-based prefilter → **1,094** gate-passing extractions

Why gate at all?
Hand-reading reviews is anecdote, not evidence — and an ungated LLM pass returns a complaints dump. The gate forces every kept review through one schema, so the output is evidence about widening baskets.

1 Ingest
15,820 REVIEWS
- ✓ Public app-store reviews
- ✓ Deduplicated before any model runs

→

2 Heuristic prefilter
→ 4,740 KEPT
- ✓ Drops empty text and reviews under five words
- ✓ Drops 5-star reviews
- ✓ Rule-based only — no LLM, no cost

→

3 Gated extraction
→ 1,094 KEPT
- ✓ Gate: is this about trying or avoiding a category?
- ✓ Extract: behaviour, category named, friction theme
- ✓ Two steps, never one

→

4 Deterministic clustering
THEMES = DB ROWS
- ✓ Frequency counted in the database
- ✓ Never an LLM estimate
- ✓ 70% land in a single theme

→

5 LLM-judge validation
HELD-OUT SAMPLE
- ✓ Gate decisions re-checked on a held-out sample
- ✓ Validation dial reported in the app itself

The engine answers the brief's discovery questions directly: why users repeat, what blocks new categories, what frustrations recur, and which unmet needs persist. The full workflow is live — open it from the link above.

✓ KEPT — CATEGORY BEHAVIOUR

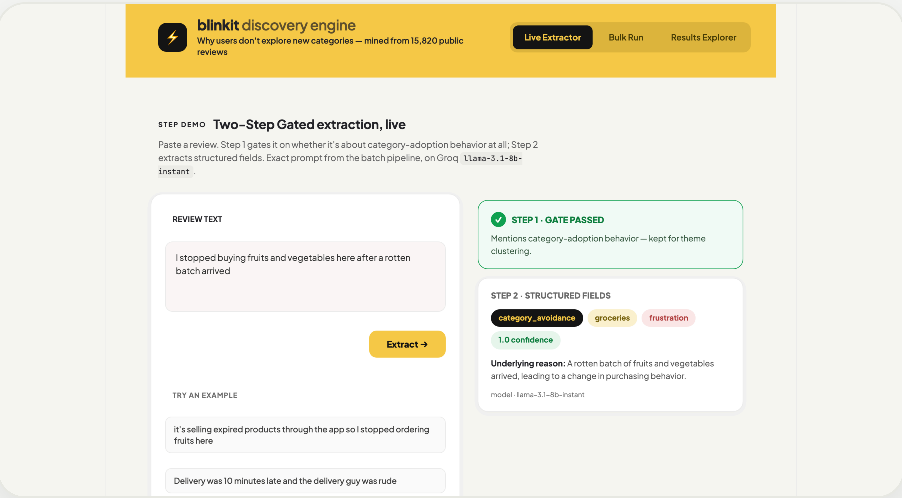

A rotten-batch review is kept, tagged category_avoidance · groceries.

✗ DROPPED — SERVICE COMPLAINT

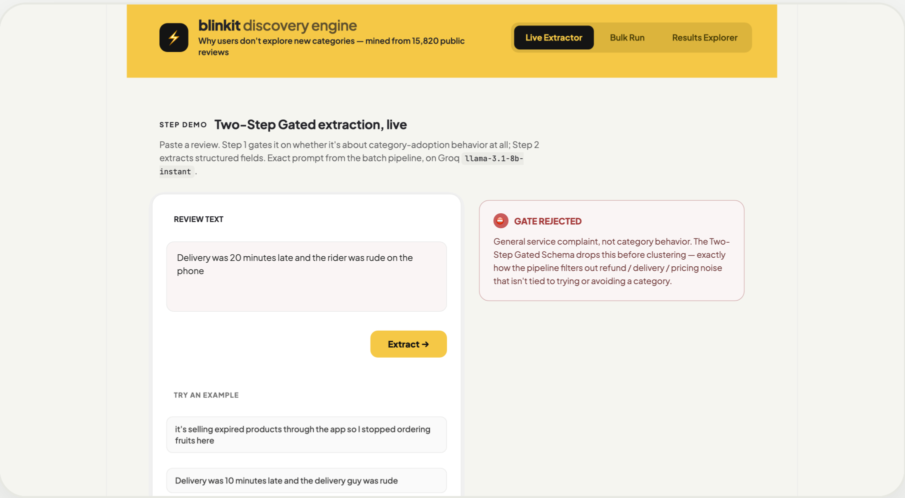

A late-delivery complaint is rejected — same engine, same review box. The gate is the difference.

! DESIGN DECISION · THE GATE

An early ungated run only surfaced generic service complaints.

Extraction now hard-gates on category-adoption relevance before anything is written to the database.

WHAT THE ENGINE MEASURES
- ✓ Adoption behaviour, per review
- ✓ The category named
- ✓ The friction theme behind it

Groq llama-3.1-8b-instant throughout.

Poor quality and unreliable products is the dominant reason customers stop widening their baskets

1 One theme carries the segment

Of the 1,094 gate-passing extractions, the dominant friction theme is “Poor Quality and Unreliable Products”.

70% of kept evidence sits in this one theme

2 The steep funnel is the gate working, not data loss

The prefilter drops empty text, reviews under five words and 5-star reviews. The gate then removes everything unrelated to trying or avoiding a category.

15,820 → 4,740 → 1,094

3 Where the top theme lands

Groceries	431
Electronics	96
Snacks & beverages	79
Household essentials	39
Personal care	27

Top five named categories inside the top theme, by extraction count.

4 Every count is a database row

Frequencies come from the pipeline export — never an LLM estimate, never a per-user confidence score. The validation dial re-checks gate decisions on a held-out sample.

THE RUN, IN THE PRODUCT

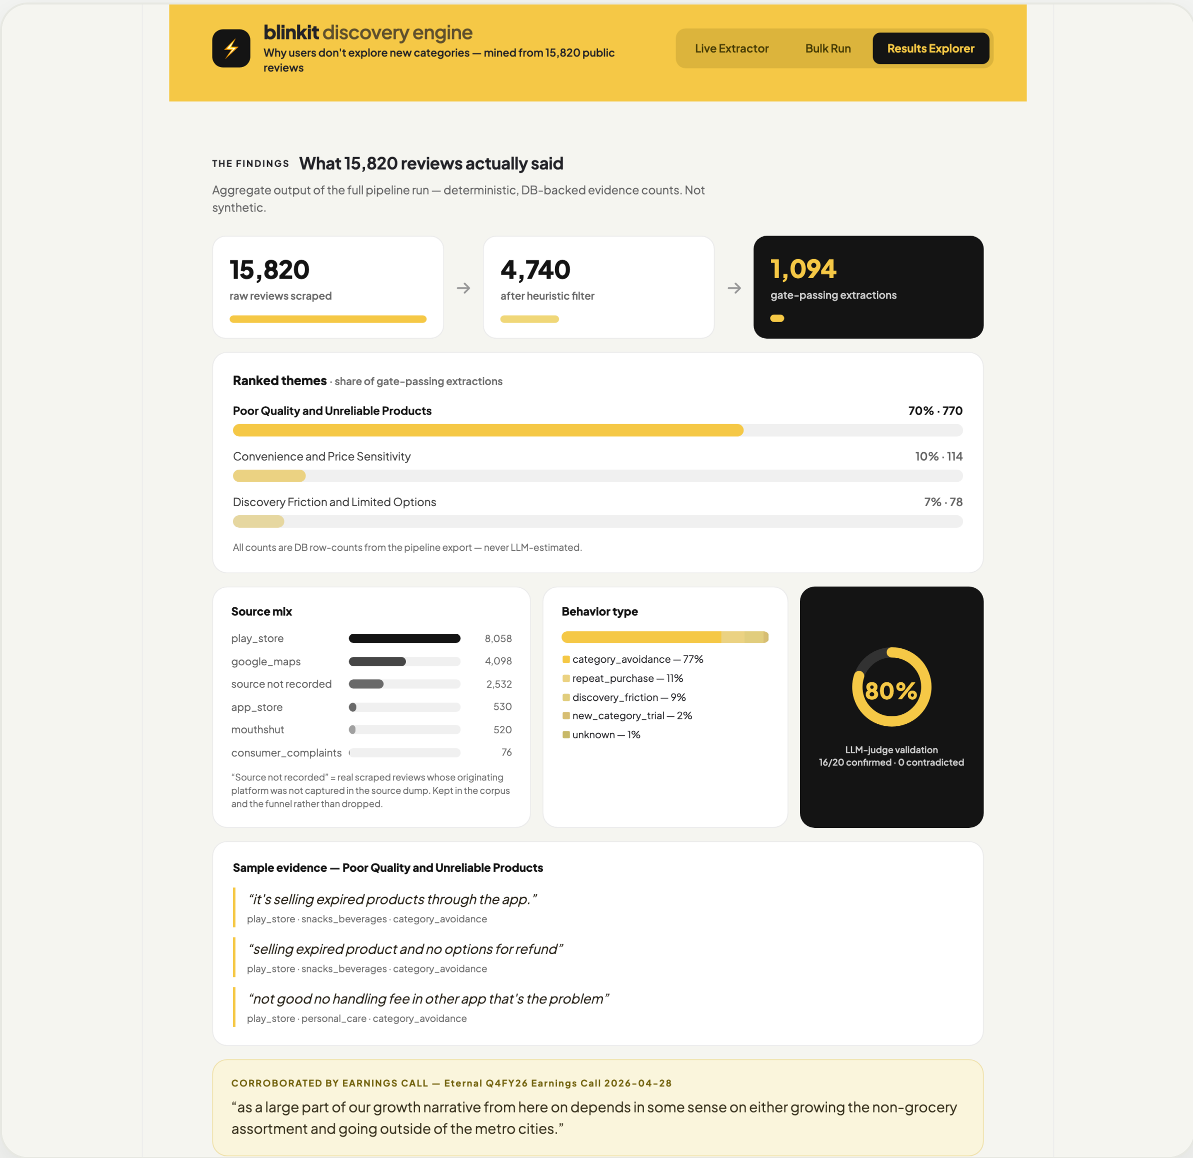

Results Explorer — funnel, ranked themes, source mix, behaviour split and the LLM-judge validation dial.

THE CORE PROBLEM Customers are not missing the other categories — they have decided those categories are unsafe to buy here. One damaged, expired or fake item is read as evidence about the whole category.

The survey confirmed the AI's cause — and challenged how far it reaches

SAMPLE, STATED HONESTLY

31 survey responses. No live interviews were conducted — a deliberate trade of depth for breadth against a fixed deadline. Six respondents consented to follow-up calls that were not run.

% Quantitative validation

- **58%** — Stick to the same categories — 18 of 31
- **48%** — Had a bad order in 3 months — 15 of 31
- **73%** — Of those, fully refunded — 11 of 15

ASKED DIRECTLY, 1–5: MEAN 3.0

▪ 11 agree ▪ 9 neutral ▪ 11 disagree

Does a bad experience in one category make you hesitant to try others?

! What the survey challenged

The behaviour change is an even split, not a rule.
Of the 15 with a bad order: 7 confirmed a change in behaviour, 7 reported no change, 1 ambiguous. The split stayed exactly even when the sample grew.

Half the "stuck" group had no incident at all.
8 of the 18 who stick to the same categories reported no bad order — "Nothing in particular — I just don't need those categories". That stagnation is low intent, not lost trust.

Some churn instead of narrowing.
"I have started using Zepto more" — platform switching rather than category avoidance.

Packaging is a distinct sub-cause.
"often packed at the bottom of the bag under heavier groceries" — a fulfilment failure, not product sourcing. The same respondent wrote in "better and safer packaging guarantees for fragile items" unprompted, rather than picking any option offered.

2 respondents raised pricing versus competitors unprompted.

✓ Confirmed, in their words

"I have stopped exploring new categories on Blinkit. I only order the essential items now if anything urgent."
Unresolved puja-item complaint

"yes I am skeptical of buying electronic products from Blinkit and other similar applications"
Damaged electronics

"I avoid buying any fresh products dairy or perishables items from quick commerce apps."
Expired dairy — generalised beyond Blinkit

The 15 bad orders: 11 damaged · 2 expired · 1 fake or duplicate · 1 other. Resolution: 11 full refunds · 3 replacements · 1 unresolved.

THE VERDICT

Confirmed: a quality failure does generalise into avoiding the whole category. **Challenged:** it does not do so for everyone — so the fix is scoped to the trust-driven share, not all stagnation.

One quality failure generalises into permanent avoidance of an entire category

01 · WHO IT'S FOR

Heavy, habitual repeat buyers.

High-frequency customers who default to the same narrow set of categories every month. 18 of 31 surveyed describe themselves this way.

02 · THE ROOT CAUSE

Trust, not discovery.

A damaged, expired or fake item teaches the customer that a category is unsafe to buy here — so they retreat to essentials, or leave for another app. It is not an awareness problem: these buyers already see the other categories every week.

03 · HOW THEY COPE TODAY

Two workarounds, both costly.

Retreat to essentials only — "I only order the essential items now if anything urgent." Or switch platforms outright — "I have started using Zepto more." The first suppresses category adoption silently; the second is active share loss.

04 · WHY IT'S WORTH IT FOR THE USER

They want to consolidate — they just can't risk it.

These customers already trust Blinkit for staples and would rather buy more in one app than juggle three. 27 of 31 named at least one concrete change that would let them — only 4 said "nothing in particular" — led by a refund guarantee (13 picks) and a visible quality signal (10). No new need is being invented.

05 · WHY IT'S WORTH IT FOR THE BUSINESS

Basket width is the growth lever the company has named.

Restoring trust re-opens non-grocery categories for customers who already order frequently — no new acquisition required.

SCOPE LIMIT · STATED UP FRONT

This addresses only the quality- and trust-driven share of stagnation.

Roughly half the stuck segment reported no incident at all. Those customers are low intent; no trust fix will move them, and we do not claim it will.

Evidence base: 1,094 gate-passing review extractions (70% quality-driven) and 31 survey responses (15 with a recent bad order) point at the same failure.

A trust-led nudge wins because it answers the blocker customers actually named

OPTIONS CONSIDERED

Option	Pain it solves	Whole segment?	Differentiated?
Discount or intro offer on a new category	Price hesitation only — 2 respondents raised pricing	✗ No — misses the trust blocker	✗ No — every rival can copy it
Better merchandising and category placement	Awareness — ranked 3rd at 7% of evidence	✗ No — these buyers already see the categories	✗ Partly only
Let users inspect before accepting delivery	Pre-purchase doubt — now the 3rd-ranked driver, 7 of 31	✗ Partly — helps at the door, not at the moment of choosing	✗ No — an ops change any rival can match, and it slows every handover
Trust-led category nudge ← chosen	**The top two named blockers: refund certainty 13/31, then quality 10/31**	✓ Yes for the trust-driven share	✓ Yes — per-user, per-friction copy

Ranked survey drivers, Q15 “what would make you try a new category” (pick up to 2), N=31: no-questions-asked refund or return guarantee 13 · better visible quality and freshness guarantees 10 · inspect before accepting delivery 7 · item reviews and ratings 5 · intro offers 5 · nothing in particular 4 · human support agent 3 · safer packaging for fragile items 1 (written in, not offered as an option).

CHOSEN SOLUTION

The Category Nudge Agent

A per-user nudge into an adjacent category that leads with the refund or return guarantee, then the quality and freshness signal — in that order, because that is the order customers ranked them.

Why not a discount

A discount pays a customer to repeat the experience that broke their trust. It does not change the expected outcome of the order.

blinkit | PART 4 · THE MVP, DEPLOYED | Open the live agent →

The agent runs today: a shopper-facing nudge, its reasoning, and the push queue behind it

DATA DISCLOSURE

8 synthetic user profiles, labelled SYNTHETIC in the product UI itself. No real customer data was available.

Leads with the two research-ranked trust drivers

No invented pricing — labelled illustrative

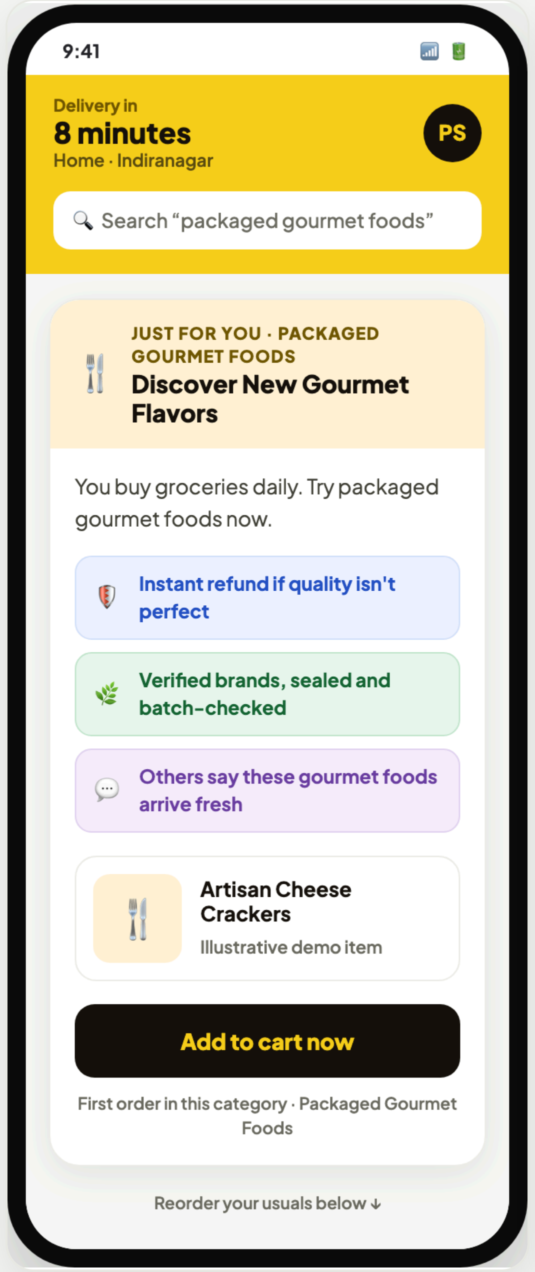

The shopper's view — a trust-led nudge into one adjacent category.

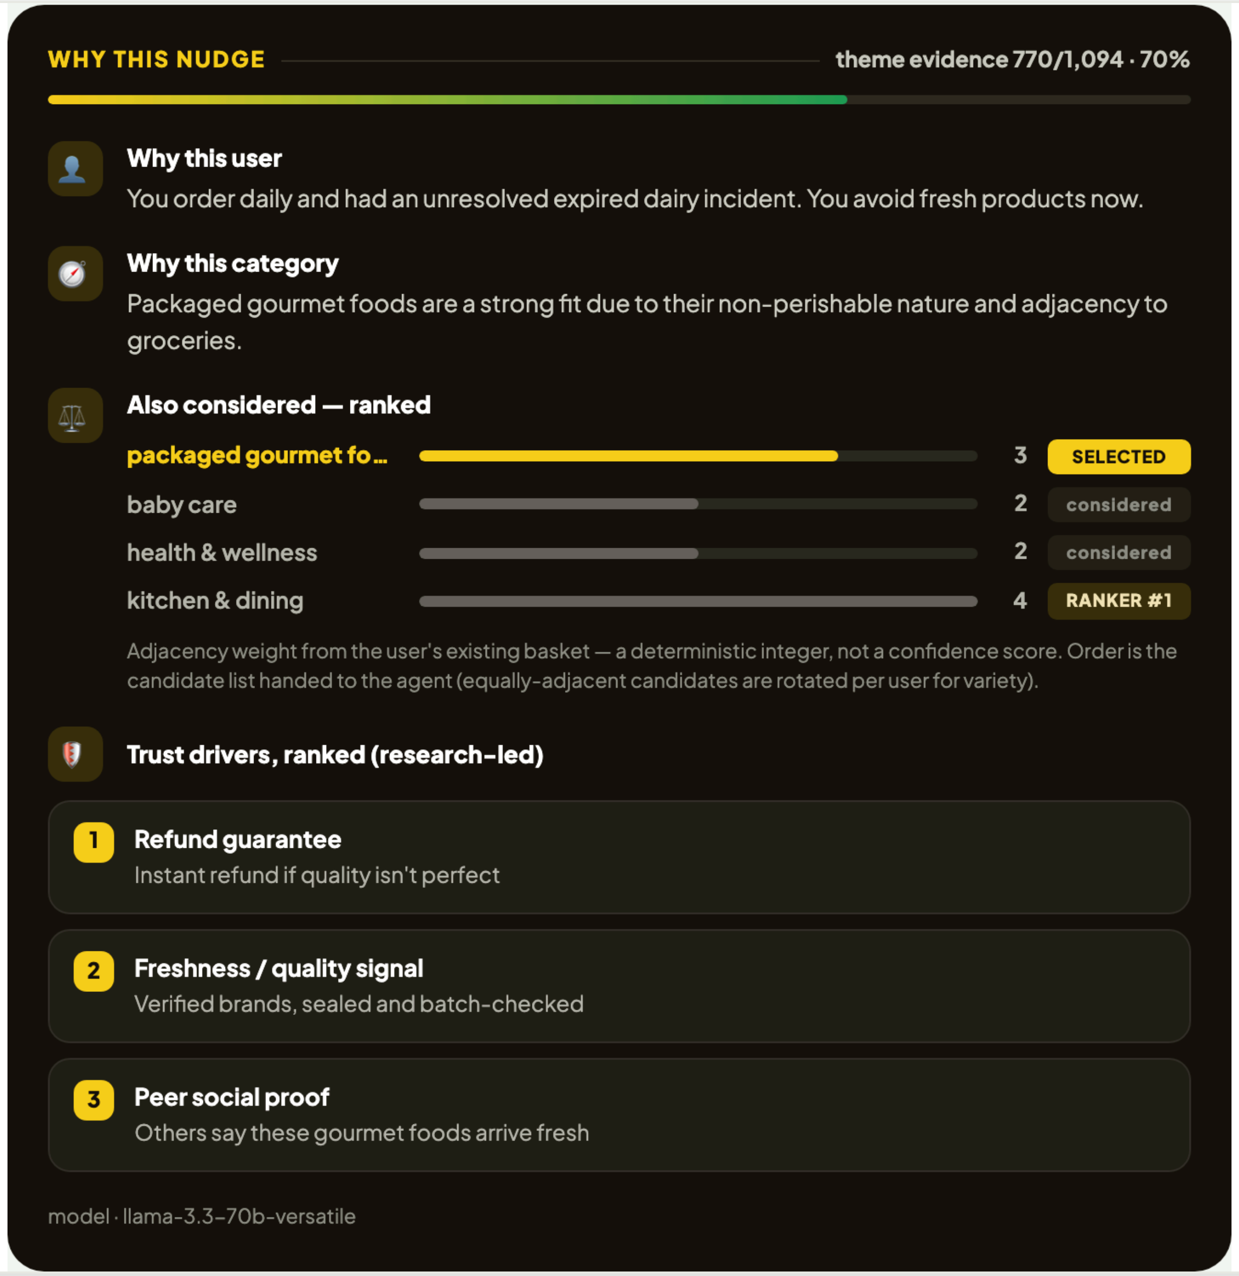

The agent's reasoning panel — why this user, why this category, and what else was considered.

Deterministic ranker, integer weights — not model confidence

Copy generated live per user, not templated

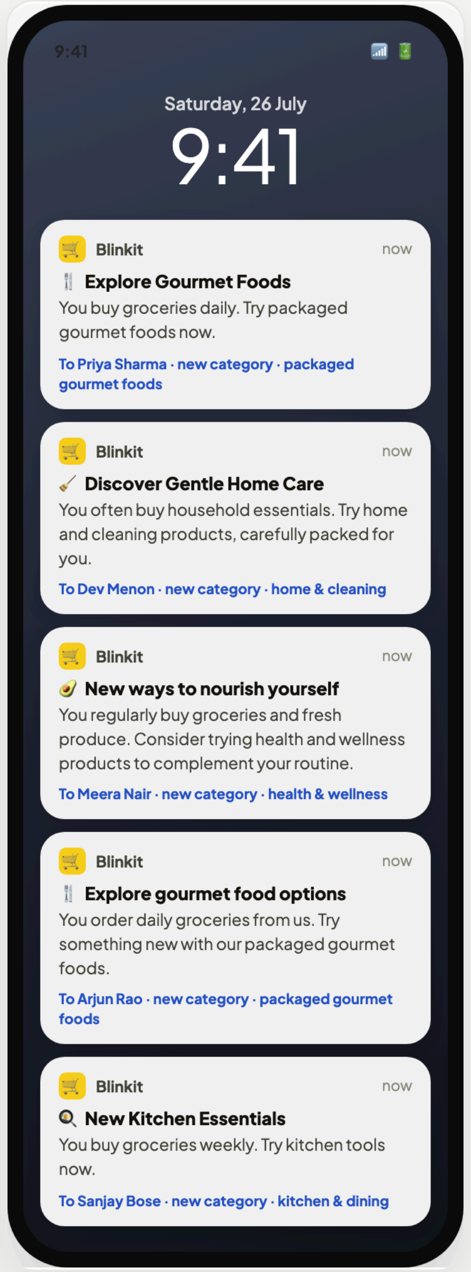

Push payloads from the auto-nudge queue.

Five users, four different categories — no collapse to one suggestion

Deterministic logic decides what to nudge; the LLM only writes the words

WHY THIS NUDGE — out of primary scope

⚠ This user shows **no quality incident** — their stagnation looks intent-driven. The agent leads with relevance instead of a guarantee, and does not claim the trust fix applies.

Why this user
You order weekly, but haven't tried new categories. You mentioned not needing them, but we think you might find something useful.

Why this category
Health and wellness is a natural fit with your existing grocery and fresh produce purchases.

Also considered — ranked

health & wellness	1	SELECTED
home & cleaning	1	considered
kitchen & dining	4	RANKER #1
packaged gourmet fo...	2	considered

Adjacency weight from the user's existing basket — a deterministic integer, not a confidence score. Order is the candidate list handed to the agent (equally-adjacent candidates are rotated per user for variety).

Trust drivers, ranked (research-led)

1. Refund guarantee
2. Freshness / quality signal

model · llama-3.3-70b-versatile

The product itself says when the fix does NOT apply

Amber banner, top of the panel — a low-intent user is told the trust fix isn't theirs.

RUNTIME FLOW

Input User profile — cadence, tenure, category history, prior incident

Rule-based Friction match — profile to a Part 1 theme, no LLM

Rule-based Adjacency ranker — integer weights produce the candidate list

LLM Nudge generation — live Groq call writes copy and reasoning

Output Rendered nudge — refund guarantee first, then quality and freshness

TWO INTEGRITY RULES, ENFORCED IN THE PROMPT

Never invent statistics. No ratings, buyer counts or success rates may appear in nudge copy — so none do.

Never overclaim to low-intent users. With no incident on file, the agent states that the trust fix does not apply.

WHERE IT BREAKS

Profiles are synthetic — nothing is matched against real behavioural history.

A refund promise is only credible if operations honour it; the nudge inherits fulfilment risk it cannot control.

Packaging failures need a fulfilment change, not a message.

A customer who has already switched platforms may never see the nudge.

blinkit | MEASUREMENT · PROPOSED, NOT YET MEASURED

We would judge this on new-category adoption, in three tiers, with trust guardrails

Plan only. No experiment has been run and the MVP has never been shipped to real users, so no result numbers appear on this slide.

PRIMARY METRIC

% of monthly active customers who purchase from at least one new category that month

Measured against the nudged segment's own prior month, not against a headline average.

HONEST SCOPE LIMIT

This moves the trust-driven share of stagnation only.

Around half the stuck segment reported no incident. Their stagnation is low intent, and a trust message is the wrong instrument for it. Any evaluation must report the two sub-segments separately or it will read as a failure of the wrong thing.

MEASUREMENT LADDER

Tier	What we would watch	Reads as
1 · Outcome	New-category purchase rate among nudged trust-blocked users	Did the behaviour change
2 · Mechanism	Nudge shown → opened → category browsed → item added	Which step leaks
3 · Durability	Repeat purchase in the newly tried category the following month	Trust restored, not just tested

GUARDRAILS

Return and refund rate in nudged categories must not rise — a nudge into a genuinely bad category is a trust loss, not a win.

Complaint rate among nudged users, tracked against their own baseline.

Order frequency and nudge fatigue — reorder rate must not fall; opt-outs capped per user.

Copy audit — sampled nudges checked for invented statistics; any hit is a release blocker.

WHAT WE'D BUILD NEXT

1 · Let users inspect before accepting

The 3rd-ranked driver (7 of 31) and the highest-ranked thing this MVP deliberately does not do. It is an operations change, not a message — which is exactly why it was scoped out.

2 · A relevance-led nudge for the low-intent half

8 of 18 stuck users had no bad experience at all. A trust message is the wrong instrument for them; the open question is whether a cart-level prompt moves them at all.

3 · Real behavioural data instead of synthetic profiles

Every profile here is synthetic. The first real test is whether the friction match holds against actual order history.